

Creando un pacman en k8s

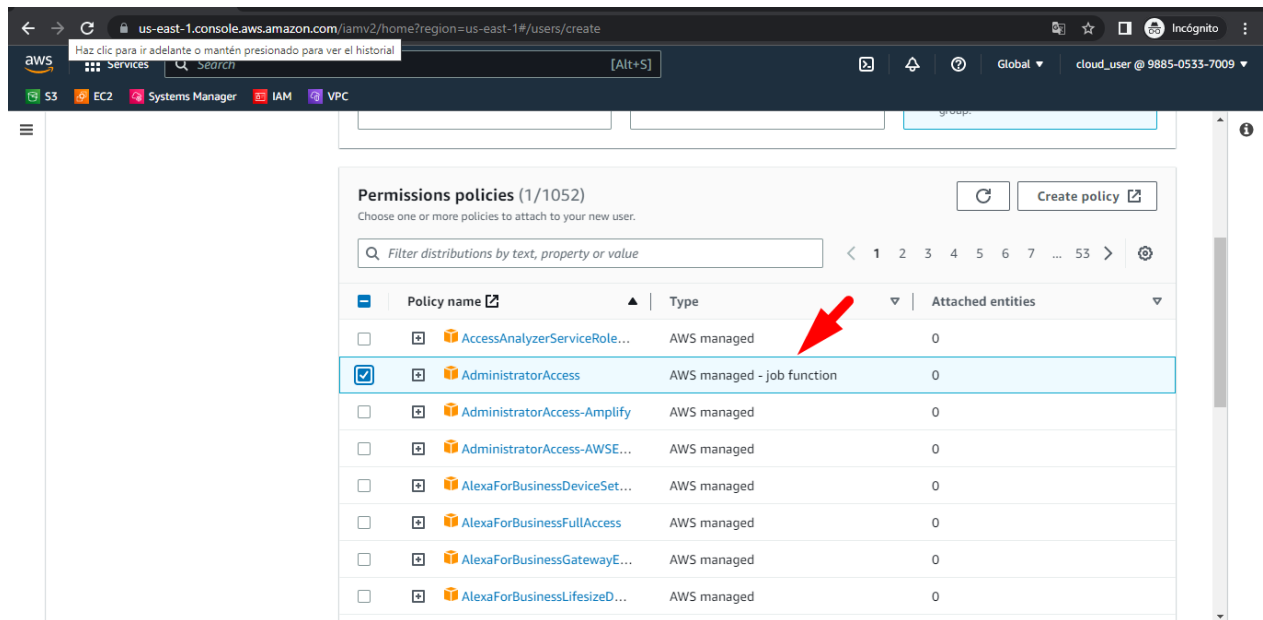
Primero creamos una cuenta de iam con acceso programatico.

The screenshot shows the AWS IAM console interface for creating a new user. The breadcrumb navigation is 'IAM > Users > Create user'. The left sidebar shows three steps: 'Step 1: Specify user details' (active), 'Step 2: Set permissions', and 'Step 3: Review and create'. The main content area is titled 'Specify user details' and contains a 'User details' section. In this section, the 'User name' field is filled with 'k8sadmin'. Below the field, a note states: 'The user name can have up to 64 characters. Valid characters: A-Z, a-z, 0-9, and + = , _ @ - (hyphen)'. There is an unchecked checkbox for 'Provide user access to the AWS Management Console - optional' with a note that it's a best practice to manage access in IAM Identity Center. A blue information box at the bottom of the section states: 'If you are creating programmatic access through access keys or service-specific credentials for AWS CodeCommit or Amazon Keyspaces, you can generate them after you create this IAM user. Learn more'. At the bottom right of the form are 'Cancel' and 'Next' buttons.

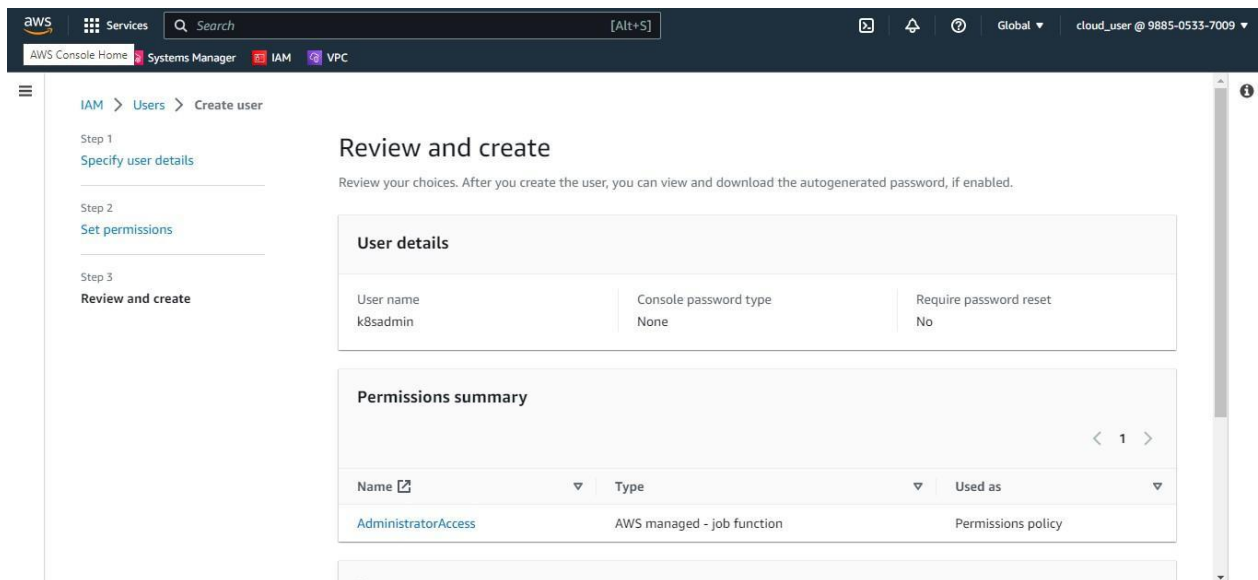
Ahora pedimos elegir una policy

The screenshot shows the 'Set permissions' step of the AWS IAM console. The breadcrumb navigation is 'IAM > Users > Create user'. The left sidebar shows three steps: 'Step 1: Specify user details', 'Step 2: Set permissions' (active), and 'Step 3: Review and create'. The main content area is titled 'Set permissions' and includes a note: 'Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. Learn more'. Below this is a 'Permissions options' section with three radio buttons: 'Add user to group' (selected), 'Copy permissions', and 'Attach policies directly'. The 'Add user to group' option has a description: 'Add user to an existing group, or create a new group. We recommend using groups to manage user permissions by job function.' The 'Copy permissions' option has a description: 'Copy all group memberships, attached managed policies, and inline policies from an existing user.' The 'Attach policies directly' option has a description: 'Attach a managed policy directly to a user. As a best practice, we recommend attaching policies to a group instead. Then, add the user to the appropriate group.' A red arrow points to the 'Attach policies directly' option. Below these options is a blue information box titled 'Get started with groups' with the text: 'Create a group and select policies to attach to the group. We recommend using groups to manage user permissions by job function, AWS service access, or custom permissions. Learn more'. A 'Create group' button is located to the right of this box.

Le damos los permisos



Ahora



Al acceso programático lo seteamos más adelante.

Luego de crear el user debería crear así....

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings

Access reports

- Access analyzer

Introducing the new Users list experience
We've redesigned the Users list experience to make it easier to use. [Let us know what you think.](#)

User created successfully
You can view and download the user's password and email instructions for signing in to the AWS Management Console. [View user](#)

Users (2) Info
An IAM user is an identity with long-term credentials that is used to interact with AWS in an account.

Find users by username or access key

	User name	Groups	Last activity	MFA	Password age	Activ
<input type="checkbox"/>	cloud_user	None	^	None	✓ 50 minutes ago	
<input type="checkbox"/>	k8sadmin	None	^	None	None	

Seleccionamos ahí para darle el acceso programático

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings

Access reports

- Access analyzer
- Archive rules

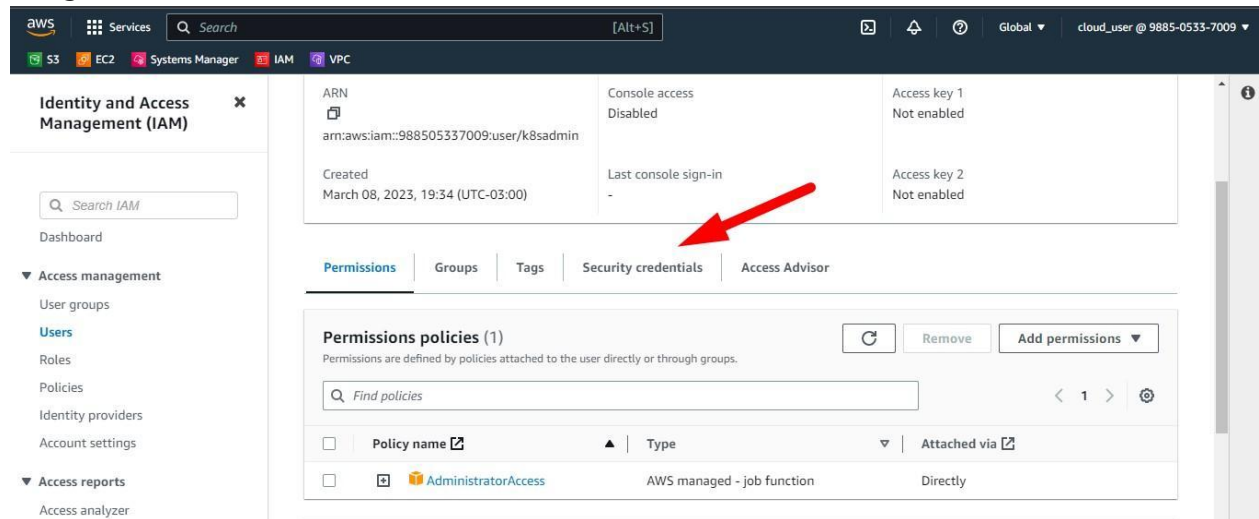
Introducing the new Users list experience
We've redesigned the Users list experience to make it easier to use. [Let us know what you think.](#)

Users (2) Info
An IAM user is an identity with long-term credentials that is used to interact with AWS in an account.

Find users by username or access key

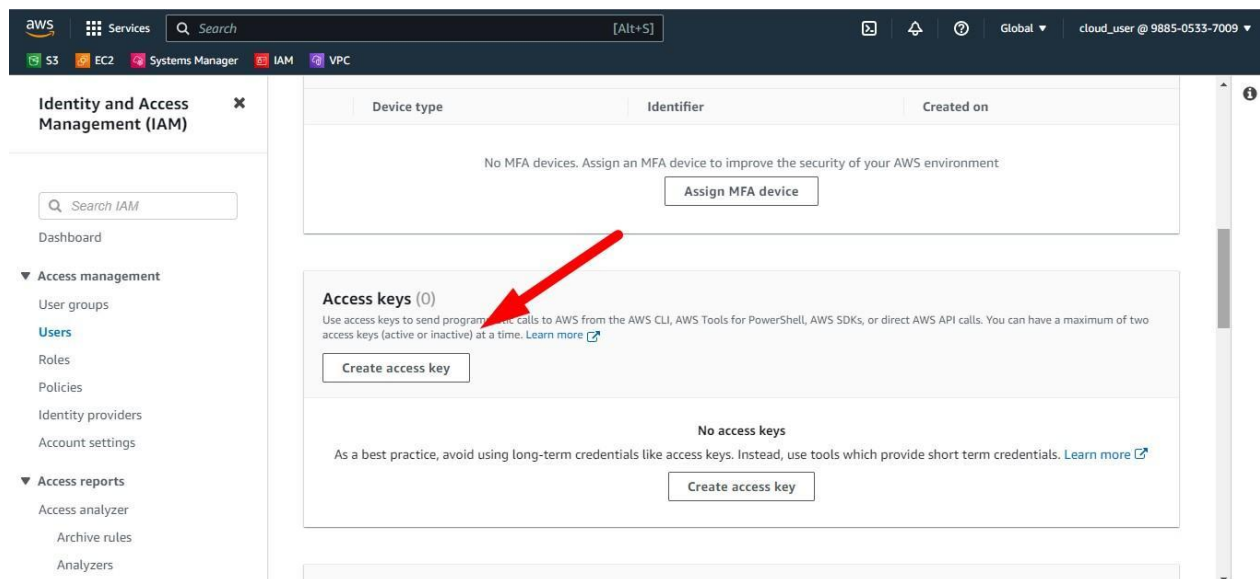
	User name	Groups	Last activity	MFA	Password age	Activ
<input type="checkbox"/>	cloud_user	None	✓ 17 minutes ago	None	✓ 52 minutes ago	✓ 51
<input type="checkbox"/>	k8sadmin	None	Never	None	None	-

Luego:



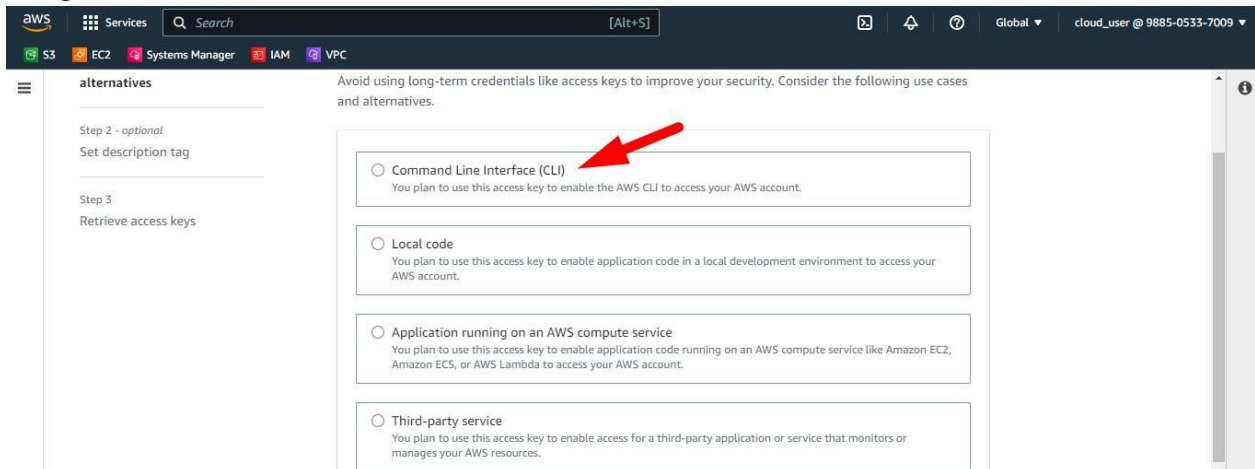
The screenshot shows the AWS IAM console interface. On the left is the 'Identity and Access Management (IAM)' sidebar with navigation links like 'Dashboard', 'Access management', 'Users', 'Roles', 'Policies', 'Identity providers', 'Account settings', 'Access reports', and 'Access analyzer'. The main content area is titled 'Security credentials' and shows details for a user with ARN 'arn:aws:iam::988505337009:user:k8sadmin'. It displays 'Console access' as 'Disabled', 'Access key 1' as 'Not enabled', and 'Access key 2' as 'Not enabled'. The 'Last console sign-in' is shown as '-'. A red arrow points to this field. Below this, the 'Permissions policies (1)' section shows a table with one policy: 'AdministratorAccess', which is 'AWS managed - job function' and 'Attached via Directly'.

Le damos en access key

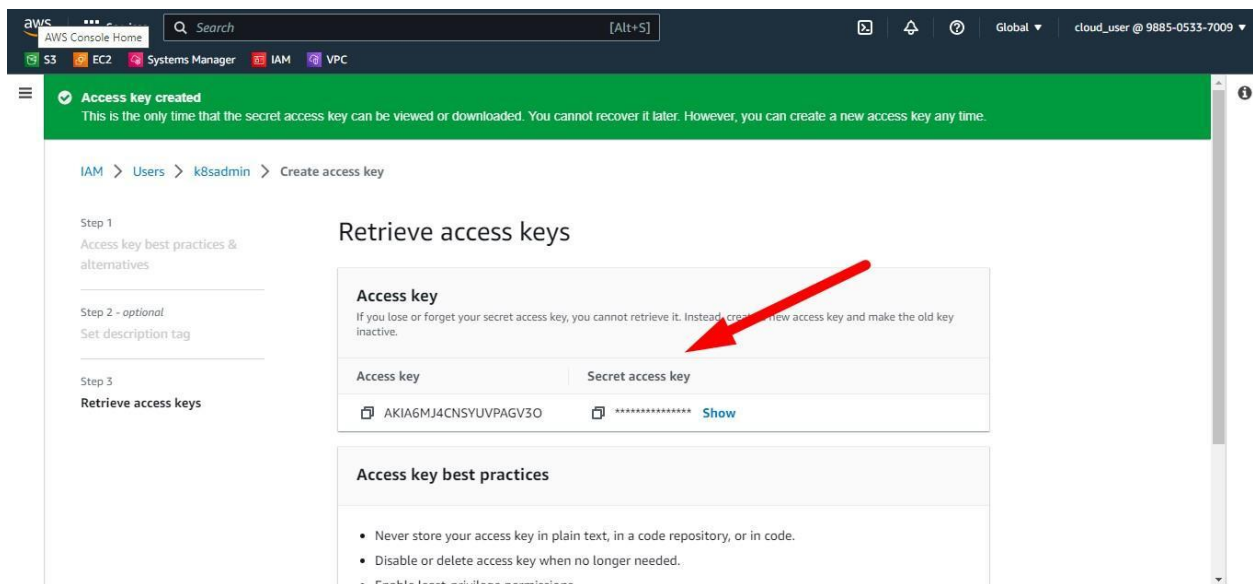


The screenshot shows the 'Access keys' section of the AWS IAM console. It displays a message: 'No MFA devices. Assign an MFA device to improve the security of your AWS environment.' with an 'Assign MFA device' button. Below this, the 'Access keys (0)' section shows a message: 'Use access keys to send programmatic calls to AWS from the AWS CLI, AWS Tools for PowerShell, AWS SDKs, or direct AWS API calls. You can have a maximum of two access keys (active or inactive) at a time. Learn more'. A red arrow points to the 'Create access key' button. At the bottom, the 'No access keys' section shows a message: 'As a best practice, avoid using long-term credentials like access keys. Instead, use tools which provide short term credentials. Learn more' with another 'Create access key' button.

Luego



aceptamos los dos checkbox y nos aparecen las credenciales



Las copiamos.

Creamos una instancia.

aws Services Search [Alt+S] N. Virginia cloud_user @ 9885-0533-7009

S3 EC2 Systems Manager IAM VPC

EC2 > Instances > Launch an instance

Launch an instance [Info](#)

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags [Info](#)

Name [Add additional tags](#)

Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below.

Summary

Number of instances [Info](#)

Software Image (AMI)
Amazon Linux 2 Kernel 5.10 AMI...[read more](#)
ami-006dcf34c09e50022

Virtual server type (instance type)
t2.micro

Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

[Cancel](#) [Launch instance](#)

Generamos el par key (bastion.pem)

aws Services Search [Alt+S] N. Virginia cloud_user @ 9885-0533-7009

S3 EC2 Systems Manager IAM VPC

Key pair name - required
Select

Please choose a key pair or choose the option to proceed without a key pair.

Create key pair

[Info](#)

We noticed that you didn't select a key pair. If you want to be able to connect to your instance it is recommended that you create one.

Key pairs allow you to connect to your instance securely.

Enter the name of the key pair below. When prompted, store the private key in a secure and accessible location on your computer. **You will need it later to connect to your instance.** [Learn more](#)

☒ Create new key pair ☐ Proceed without key pair

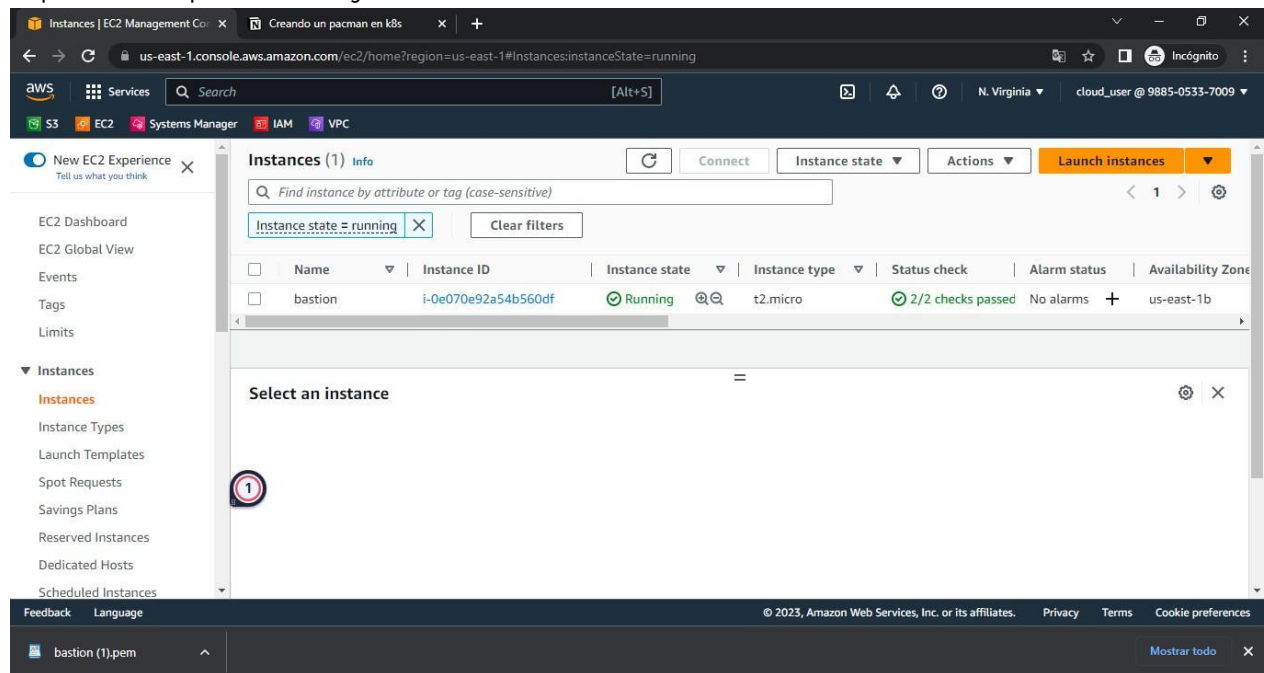
Key pair name

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type
☒ RSA
RSA encrypted private and public key pair
☐ ED25519

[Launch instance](#)

Esperamos que se cree y entremos.



Ahora nos logueamos con la clave generada

```
ssh -i bastion.pem ec2-user@54.90.140.133
Warning: Permanently added '54.90.140.133' (RSA) to the list of known hosts.
X11 forwarding request failed on channel 0

_ _ _ _ _
_ _ | < / Amazon Linux 2 AMI
_ _ | \ _ _ |

https://aws.amazon.com/amazon-linux-2/
4 package(s) needed for security, out of 6
available Run "sudo yum update" to apply all
updates.
```

y empezamos a instalar las tooles que necesitamos....

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip" && unzip awscliv2.zip
```

Luego

```
sudo ./aws/install --bin-dir /usr/bin --install-dir /usr/bin/aws-cli --update
aws --version
```

Luego

Configuramos aws con la data programatica

```
aws
configure
```

Ahora bajamos kubectl

```
curl -o kubectl https://amazon-eks.s3.us-west-2.amazonaws.com/1.16.8/2020-04-16/bin/linux/amd64/kubectl
```

Lo seteamos

```
chmod +x ./kubectl
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export PATH=$PATH:$HOME/bin
kubectl version
```

Bajamos las configs que necesitamos

```
wget https://github.com/pluralsight-cloud/content-deploying-and-managing-a-web-application-in-kubernetes-with-terraform/raw/main/eks.zip
```

descomprimos

```
[ec2-user@ip-172-31-30-154 ~]$ unzip
eks.zip Archive: eks.zip
  creating: eks/
  inflating: eks/
  eks/eks-cluster.tf
  inflating: eks/main.tf
  inflating: eks/outputs.tf
  inflating: eks/terraform.tf
  inflating: eks/variables.tf
  inflating: eks/vpc.tf
```

Luego instalar el cliente terrarform

```
sudo yum install -y yum-utils
sudo yum-config-manager --add-repo https://rpm.releases.hashicorp.com/AmazonLinux/hashicorp.repo
sudo yum -y install terraform git
```

y finalmente la secuencia de siempre....


```

joe2-user@ip-172-31-38-354 ~$ kubectl kubernetes
kubectl initializing modules...
Downloading registry.kubernetes.io/kubernetes-aws-modules/eks/aws 19.8.4 for eks...

- eks in .kubernetes/modules/eks
- eks_eks_managed_node_group in .kubernetes/modules/eks/modules/eks-managed-node-group
- eks_eks_managed_node_group_aws_data in .kubernetes/modules/eks/modules/_aws_data
- eks_fargate_profile in .kubernetes/modules/eks/modules/fargateprofile
Downloading registry.kubernetes.io/kubernetes-aws-modules/eks/aws 1.1.0 for eks_aws...
- eks_aws in .kubernetes/modules/eks_aws
- eks_aws_managed_node_group in .kubernetes/modules/eks/modules/eks-managed-node-group
- eks_aws_managed_node_group_aws_data in .kubernetes/modules/eks/modules/_aws_data
Downloading registry.kubernetes.io/kubernetes-aws-modules/eks/aws 3.14.2 for eks_aws...
- eks_aws in .kubernetes/modules/eks_aws

Initializing the
kubernetes...
Installing provider plugins...
- Finding hashicorp/kubesection versions matching "== 2.10.0, != 2.10.3"...
- Finding hashicorp/aws versions matching "== 3.40.0, != 3.42.0, != 3.43.0, != 3.44.0"...
- Finding hashicorp/random versions matching ">= 2.4.0"...
- Finding hashicorp/tls versions matching "== 3.0.0, != 4.0.4"...
- Finding hashicorp/consul versions matching "== 3.0.0, != 3.0.0"...
- Installing hashicorp/kubesection v2.10.0...
- Installed hashicorp/kubesection v2.10.0 (signed by HashiCorp)
- Installing hashicorp/aws v3.40.0...

```

```

- Installing hashicorp/cloudinit v2.2.0...
- Installed hashicorp/cloudinit v2.2.0 (signed by HashiCorp)

Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control
repository so that Terraform can guarantee to make the same selections by
default when you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to
see any changes that are required for your infrastructure. All Terraform
commands should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget,
other commands will detect it and remind you to do so if necessary.
[ec2-user@ip-172-31-30-154 eks]$

```

Lanzamos el terraform plan

```

[ec2-user@ip-172-31-30-154 eks]$ terraform plan data.aws_availability_zones.available:
Reading... module.eks.module.eks_managed_node_group["one"].data.aws_partition.current:
Reading... module.eks.data.aws_partition.current: Reading...
module.eks.module.eks_managed_node_group["two"].data.aws_partition.current: Reading...
module.eks.module.eks_managed_node_group["two"].data.aws_caller_identity.current: Reading...
module.eks.module.kms.data.aws_caller_identity.current: Reading...
module.eks.module.kms.data.aws_partition.current: Reading...
module.eks.data.aws_caller_identity.current: Reading... module.eks.data.aws_partition.current:
Read complete after 0s [id=aws]
module.eks.module.eks_managed_node_group["two"].data.aws_partition.current: Read complete after 0s [id=aws]
module.eks.module.eks_managed_node_group["one"].data.aws_partition.current: Read complete after 0s [id=aws]
module.eks.module.kms.data.aws_partition.current: Read complete after 0s [id=aws]
module.eks.module.eks_managed_node_group["two"].data.aws_caller_identity.current: Read complete after 0s [id=988505337009]
module.eks.module.eks_managed_node_group["one"].data.aws_iam_policy_document.assume_role_policy[0]: Reading...
module.eks.module.eks_managed_node_group["one"].data.aws_caller_identity.current: Reading...
module.eks.module.eks_managed_node_group["two"].data.aws_iam_policy_document.assume_role_policy[0]: Reading...
module.eks.data.aws_iam_policy_document.assume_role_policy[0]: Reading...
module.eks.data.aws_caller_identity.current: Read complete after 0s [id=988505337009] module.eks.module.kms.data.aws_caller_identity.current: Read
complete after 0s [id=988505337009] module.eks.module.eks_managed_node_group["one"].data.aws_caller_identity.current: Read complete after 0s
[id=988505337009] module.eks.data.aws_iam_policy_document.assume_role_policy[0]: Read complete after 0s [id=2764486067]
module.eks.module.eks_managed_node_group["one"].data.aws_iam_policy_document.assume_role_policy[0]: Read complete after 0s [id=2560088296]
module.eks.module.eks_managed_node_group["two"].data.aws_iam_policy_document.assume_role_policy[0]: Read complete after 0s [id=2560088296]
data.aws_availability_zones.available: Read complete after 0s [id=us-east-1]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the
following symbols:
+ create
<= read (data resources)
Terraform will perform the following actions: #
random_string.suffix will be created
+ resource "random_string" "suffix" {
  + id           = (known after apply)
  + length      = 8
  + lower       = true
  + min_lower   = 0
  + min_numeric = 0
  + min_special = 0
  + min_upper   = 0
  + number      = true
  + numeric     = true
  + result      = (known after apply)
  + special     = false
  + upper       = true
}

# module.eks.data.tls_certificate.this[0] will be read during apply #
(config refers to values not yet known)
<= data "tls_certificate" "this" {
  + certificates = (known after apply)
  + id          = (known after apply)
  + url         = (known after apply)
}

```

```

# module.eks.aws_cloudwatch_log_group.this[0] will be created
+ resource "aws_cloudwatch_log_group" "this" {
  + arn                = (known after apply)
  + id                 = (known after apply)
  + name               = (known after apply)
  + name_prefix        = (known after apply)
  + retention_in_days  = 90
  + skip_destroy       = false
  + tags_all           = (known after apply)
}

# module.eks.aws_eks_cluster.this[0] will be created
+ resource "aws_eks_cluster" "this" {
  + arn                = (known after apply)
  + certificate_authority = (known after apply)
  + cluster_id         = (known after apply)
  + created_at          = (known after apply)
  + enabled_cluster_log_types = [
    + "api",
    + "audit",
    + "authenticator",
  ]
  + endpoint            = (known after apply)
  + id                  = (known after apply)
  + identity             = (known after apply)
  + name                = (known after apply)
  + platform_version    = (known after apply)
  + role_arn             = (known after apply)
  + status               = (known after apply)
  + tags_all            = (known after apply)
  + version              = "1.24"

  + encryption_config {
    + resources = [
      + "secrets",
    ]

    + provider {
      + key_arn = (known after apply)
    }
  }

  + kubernetes_network_config {
    + ip_family      = (known after apply)
    + service_ipv4_cidr = (known after apply)
    + service_ipv6_cidr = (known after apply)
  }

  + timeouts {}

  + vpc_config {
    + cluster_security_group_id = (known after apply)
    + endpoint_private_access   = true
    + endpoint_public_access    = true
    + public_access_cidrs      = [
      + "0.0.0.0/0",
    ]
    + security_group_ids       = (known after apply)
    + subnet_ids               = (known after apply)
    + vpc_id                   = (known after apply)
  }
}

# module.eks.aws_iam_openid_connect_provider.oidc_provider[0] will be created
+ resource "aws_iam_openid_connect_provider" "oidc_provider" {
  + arn                = (known after apply)
  + client_id_list     = [
    + "sts.amazonaws.com",
  ]
  + id                 = (known after apply)
  + tags               = (known after apply)
  + tags_all           = (known after apply)
  + thumbprint_list    = (known after apply)
  + url                 = (known after apply)
}

# module.eks.aws_iam_policy.cluster_encryption[0] will be created
+ resource "aws_iam_policy" "cluster_encryption" {
  + arn                = (known after apply)
  + description        = "Cluster encryption policy to allow cluster role to utilize CMK provided"
  + id                 = (known after apply)

```

```

+ name           = (known after apply)
+ name_prefix    = (known after apply)
+ path           = "/"
+ policy         = (known after apply)
+ policy_id      = (known after apply)
+ tags_all       = (known after apply)
}

# module.eks.aws_iam_role.this[0] will be created
+ resource "aws_iam_role" "this" {
  + arn              = (known after apply)
  + assume_role_policy = jsonencode(
    {
      + Statement = [
        + {
          + Action   = "sts:AssumeRole"
          + Effect    = "Allow"
          + Principal = {
            + Service = "eks.amazonaws.com"
          }
          + Sid      = "EKSClusterAssumeRole"
        },
      ]
      + Version = "2012-10-17"
    }
  )
  + create_date      = (known after apply)
  + force_detach_policies = true
  + id               = (known after apply)
  + managed_policy_arns = (known after apply)
  + max_session_duration = 3600
  + name             = (known after apply)
  + name_prefix      = (known after apply)
  + path             = "/"
  + tags_all         = (known after apply)
  + unique_id        = (known after apply)

  + inline_policy {
    + name = (known after apply)
    + policy = jsonencode(
      {
        + Statement = [
          + {
            + Action = [
              + "logs:CreateLogGroup",
            ]
            + Effect = "Deny"
            + Resource = "*"
          },
        ]
        + Version = "2012-10-17"
      }
    )
  }
}

# module.eks.aws_iam_role_policy_attachment.cluster_encryption[0] will be created
+ resource "aws_iam_role_policy_attachment" "cluster_encryption" {
  + id           = (known after apply)
  + policy_arn   = (known after apply)
  + role         = (known after apply)
}

# module.eks.aws_iam_role_policy_attachment.this["AmazonEKSClusterPolicy"] will be created
+ resource "aws_iam_role_policy_attachment" "this" {
  + id           = (known after apply)
  + policy_arn   = "arn:aws:iam::aws:policy/AmazonEKSClusterPolicy"
  + role         = (known after apply)
}

# module.eks.aws_iam_role_policy_attachment.this["AmazonEKSVPResourceController"] will be created
+ resource "aws_iam_role_policy_attachment" "this" {
  + id           = (known after apply)
  + policy_arn   = "arn:aws:iam::aws:policy/AmazonEKSVPResourceController"
  + role         = (known after apply)
}

# module.eks.aws_security_group.cluster[0] will be created
+ resource "aws_security_group" "cluster" {
  + arn              = (known after apply)
  + description      = "EKS cluster security group"

```

```

+ egress          = (known after apply)
+ id              = (known after apply)
+ ingress         = (known after apply)
+ name            = (known after apply)
+ name_prefix     = (known after apply)
+ owner_id        = (known after apply)
+ revoke_rules_on_delete = false
+ tags            = (known after apply)
+ tags_all        = (known after apply)
+ vpc_id          = (known after apply)
}

# module.eks.aws_security_group.node[0] will be created
+ resource "aws_security_group" "node" {
+   arn              = (known after apply)
+   description      = "EKS node shared security group"
+   egress           = (known after apply)
+   id               = (known after apply)
+   ingress          = (known after apply)
+   name             = (known after apply)
+   name_prefix      = (known after apply)
+   owner_id         = (known after apply)
+   revoke_rules_on_delete = false
+   tags             = (known after apply)
+   tags_all         = (known after apply)
+   vpc_id           = (known after apply)
}

# module.eks.aws_security_group_rule.cluster["ingress_nodes_443"] will be created
+ resource "aws_security_group_rule" "cluster" {
+   description      = "Node groups to cluster API"
+   from_port        = 443
+   id               = (known after apply)
+   protocol         = "tcp"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = false
+   source_security_group_id = (known after apply)
+   to_port          = 443
+   type             = "ingress"
}

# module.eks.aws_security_group_rule.node["egress_all"] will be created
+ resource "aws_security_group_rule" "node" {
+   cidr_blocks      = [
+     + "0.0.0.0/0",
+   ]
+   description      = "Allow all egress"
+   from_port        = 0
+   id               = (known after apply)
+   prefix_list_ids   = []
+   protocol         = "-1"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = false
+   source_security_group_id = (known after apply)
+   to_port          = 0
+   type             = "egress"
}

# module.eks.aws_security_group_rule.node["ingress_cluster_443"] will be created
+ resource "aws_security_group_rule" "node" {
+   description      = "Cluster API to node groups"
+   from_port        = 443
+   id               = (known after apply)
+   prefix_list_ids   = []
+   protocol         = "tcp"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = false
+   source_security_group_id = (known after apply)
+   to_port          = 443
+   type             = "ingress"
}

# module.eks.aws_security_group_rule.node["ingress_cluster_8443_webhook"] will be created
+ resource "aws_security_group_rule" "node" {
+   description      = "Cluster API to node 8443/tcp webhook"
+   from_port        = 8443
+   id               = (known after apply)
+   prefix_list_ids   = []

```

```

+ protocol           = "tcp"
+ security_group_id   = (known after apply)
+ security_group_rule_id = (known after apply)
+ self               = false
+ source_security_group_id = (known after apply)
+ to_port             = 8443
+ type               = "ingress"
}

# module.eks.aws_security_group_rule.node["ingress_cluster_9443_webhook"] will be created
+ resource "aws_security_group_rule" "node" {
+   description       = "Cluster API to node 9443/tcp webhook"
+   from_port         = 9443
+   id               = (known after apply)
+   prefix_list_ids   = []
+   protocol          = "tcp"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = false
+   source_security_group_id = (known after apply)
+   to_port           = 9443
+   type             = "ingress"
}

# module.eks.aws_security_group_rule.node["ingress_cluster_kubelet"] will be created
+ resource "aws_security_group_rule" "node" {
+   description       = "Cluster API to node kubelets"
+   from_port         = 10250
+   id               = (known after apply)
+   prefix_list_ids   = []
+   protocol          = "tcp"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = false
+   source_security_group_id = (known after apply)
+   to_port           = 10250
+   type             = "ingress"
}

# module.eks.aws_security_group_rule.node["ingress_nodes_ephemeral"] will be created
+ resource "aws_security_group_rule" "node" {
+   description       = "Node to node ingress on ephemeral ports"
+   from_port         = 1025
+   id               = (known after apply)
+   prefix_list_ids   = []
+   protocol          = "tcp"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = true
+   source_security_group_id = (known after apply)
+   to_port           = 65535
+   type             = "ingress"
}

# module.eks.aws_security_group_rule.node["ingress_self_coredns_tcp"] will be created
+ resource "aws_security_group_rule" "node" {
+   description       = "Node to node CoreDNS"
+   from_port         = 53
+   id               = (known after apply)
+   prefix_list_ids   = []
+   protocol          = "tcp"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = true
+   source_security_group_id = (known after apply)
+   to_port           = 53
+   type             = "ingress"
}

# module.eks.aws_security_group_rule.node["ingress_self_coredns_udp"] will be created
+ resource "aws_security_group_rule" "node" {
+   description       = "Node to node CoreDNS UDP"
+   from_port         = 53
+   id               = (known after apply)
+   prefix_list_ids   = []
+   protocol          = "udp"
+   security_group_id = (known after apply)
+   security_group_rule_id = (known after apply)
+   self             = true
+   source_security_group_id = (known after apply)
+   to_port           = 53

```

```

+ type = "ingress"
}

# module.vpc.aws_eip.nat[0] will be created
+ resource "aws_eip" "nat" {
+   allocation_id = (known after apply)
+   association_id = (known after apply)
+   carrier_ip    = (known after apply)
+   customer_owned_ip = (known after apply)
+   domain        = (known after apply)
+   id            = (known after apply)
+   instance      = (known after apply)
+   network_border_group = (known after apply)
+   network_interface = (known after apply)
+   private_dns      = (known after apply)
+   private_ip       = (known after apply)
+   public_dns       = (known after apply)
+   public_ip        = (known after apply)
+   public_ipv4_pool = (known after apply)
+   tags             = {
+     + "Name" = "guru-vpc-us-east-la"
+   }
+   tags_all = {
+     + "Name" = "guru-vpc-us-east-la"
+   }
+   vpc = true
}

# module.vpc.aws_internet_gateway.this[0] will be created
+ resource "aws_internet_gateway" "this" {
+   arn = (known after apply)
+   id  = (known after apply)
+   owner_id = (known after apply)
+   tags = {
+     + "Name" = "guru-vpc"
+   }
+   tags_all = {
+     + "Name" = "guru-vpc"
+   }
+   vpc_id = (known after apply)
}

# module.vpc.aws_nat_gateway.this[0] will be created
+ resource "aws_nat_gateway" "this" {
+   allocation_id = (known after apply)
+   connectivity_type = "public"
+   id            = (known after apply)
+   network_interface_id = (known after apply)
+   private_ip      = (known after apply)
+   public_ip       = (known after apply)
+   subnet_id       = (known after apply)
+   tags            = {
+     + "Name" = "guru-vpc-us-east-la"
+   }
+   tags_all = {
+     + "Name" = "guru-vpc-us-east-la"
+   }
}

# module.vpc.aws_route.private_nat_gateway[0] will be created
+ resource "aws_route" "private_nat_gateway" {
+   destination_cidr_block = "0.0.0.0/0"
+   id                    = (known after apply)
+   instance_id           = (known after apply)
+   instance_owner_id     = (known after apply)
+   nat_gateway_id        = (known after apply)
+   network_interface_id = (known after apply)
+   origin                 = (known after apply)
+   route_table_id        = (known after apply)
+   state                 = (known after apply)

+   timeouts {
+     + create = "5m"
+   }
}

# module.vpc.aws_route.public_internet_gateway[0] will be created
+ resource "aws_route" "public_internet_gateway" {
+   destination_cidr_block = "0.0.0.0/0"
+   gateway_id            = (known after apply)
+   id                    = (known after apply)

```

```

+ instance_id           = (known after apply)
+ instance_owner_id     = (known after apply)
+ network_interface_id = (known after apply)
+ origin                = (known after apply)
+ route_table_id        = (known after apply)
+ state                 = (known after apply)

+ timeouts {
+   create = "5m"
+ }
}

# module.vpc.aws_route_table.private[0] will be created
+ resource "aws_route_table" "private" {
+   arn                = (known after apply)
+   id                 = (known after apply)
+   owner_id           = (known after apply)
+   propagating_vgws   = (known after apply)
+   route              = (known after apply)
+   tags               = {
+     "Name" = "guru-vpc-private"
+   }
+   tags_all           = {
+     "Name" = "guru-vpc-private"
+   }
+   vpc_id             = (known after apply)
+ }

# module.vpc.aws_route_table.public[0] will be created
+ resource "aws_route_table" "public" {
+   arn                = (known after apply)
+   id                 = (known after apply)
+   owner_id           = (known after apply)
+   propagating_vgws   = (known after apply)
+   route              = (known after apply)
+   tags               = {
+     "Name" = "guru-vpc-public"
+   }
+   tags_all           = {
+     "Name" = "guru-vpc-public"
+   }
+   vpc_id             = (known after apply)
+ }

# module.vpc.aws_route_table_association.private[0] will be created
+ resource "aws_route_table_association" "private" {
+   id                 = (known after apply)
+   route_table_id     = (known after apply)
+   subnet_id          = (known after apply)
+ }

# module.vpc.aws_route_table_association.private[1] will be created
+ resource "aws_route_table_association" "private" {
+   id                 = (known after apply)
+   route_table_id     = (known after apply)
+   subnet_id          = (known after apply)
+ }

# module.vpc.aws_route_table_association.private[2] will be created
+ resource "aws_route_table_association" "private" {
+   id                 = (known after apply)
+   route_table_id     = (known after apply)
+   subnet_id          = (known after apply)
+ }

# module.vpc.aws_route_table_association.public[0] will be created
+ resource "aws_route_table_association" "public" {
+   id                 = (known after apply)
+   route_table_id     = (known after apply)
+   subnet_id          = (known after apply)
+ }

# module.vpc.aws_route_table_association.public[1] will be created
+ resource "aws_route_table_association" "public" {
+   id                 = (known after apply)
+   route_table_id     = (known after apply)
+   subnet_id          = (known after apply)
+ }

# module.vpc.aws_route_table_association.public[2] will be created
+ resource "aws_route_table_association" "public" {

```



```

+ id = (known after apply)
+ route_table_id = (known after apply)
+ subnet_id = (known after apply)
}

# module.vpc.aws_subnet.private[0] will be created
+ resource "aws_subnet" "private" {
+   arn = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone = "us-east-1a"
+   availability_zone_id = (known after apply)
+   cidr_block = "10.0.1.0/24"
+   enable_dns64 = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native = false
+   map_public_ip_on_launch = false
+   owner_id = (known after apply)
+   private_dns_hostname_type_on_launch = (known after apply)
+   tags = (known after apply)
+   tags_all = (known after apply)
+   vpc_id = (known after apply)
}

# module.vpc.aws_subnet.private[1] will be created
+ resource "aws_subnet" "private" {
+   arn = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone = "us-east-1b"
+   availability_zone_id = (known after apply)
+   cidr_block = "10.0.2.0/24"
+   enable_dns64 = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native = false
+   map_public_ip_on_launch = false
+   owner_id = (known after apply)
+   private_dns_hostname_type_on_launch = (known after apply)
+   tags = (known after apply)
+   tags_all = (known after apply)
+   vpc_id = (known after apply)
}

# module.vpc.aws_subnet.private[2] will be created
+ resource "aws_subnet" "private" {
+   arn = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone = "us-east-1c"
+   availability_zone_id = (known after apply)
+   cidr_block = "10.0.3.0/24"
+   enable_dns64 = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native = false
+   map_public_ip_on_launch = false
+   owner_id = (known after apply)
+   private_dns_hostname_type_on_launch = (known after apply)
+   tags = (known after apply)
+   tags_all = (known after apply)
+   vpc_id = (known after apply)
}

# module.vpc.aws_subnet.public[0] will be created
+ resource "aws_subnet" "public" {
+   arn = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone = "us-east-1a"
+   availability_zone_id = (known after apply)
+   cidr_block = "10.0.4.0/24"
+   enable_dns64 = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native = false

```

```

+ map_public_ip_on_launch = true
+ owner_id = (known after apply)
+ private_dns_hostname_type_on_launch = (known after apply)
+ tags = (known after apply)
+ tags_all = (known after apply)
+ vpc_id = (known after apply)
}

# module.vpc.aws_subnet.public[1] will be created
+ resource "aws_subnet" "public" {
+   arn = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone = "us-east-1b"
+   availability_zone_id = (known after apply)
+   cidr_block = "10.0.5.0/24"
+   enable_dns64 = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native = false
+   map_public_ip_on_launch = true
+   owner_id = (known after apply)
+   private_dns_hostname_type_on_launch = (known after apply)
+   tags = (known after apply)
+   tags_all = (known after apply)
+   vpc_id = (known after apply)
}

# module.vpc.aws_subnet.public[2] will be created
+ resource "aws_subnet" "public" {
+   arn = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone = "us-east-1c"
+   availability_zone_id = (known after apply)
+   cidr_block = "10.0.6.0/24"
+   enable_dns64 = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native = false
+   map_public_ip_on_launch = true
+   owner_id = (known after apply)
+   private_dns_hostname_type_on_launch = (known after apply)
+   tags = (known after apply)
+   tags_all = (known after apply)
+   vpc_id = (known after apply)
}

# module.vpc.aws_vpc.this[0] will be created
+ resource "aws_vpc" "this" {
+   arn = (known after apply)
+   assign_generated_ipv6_cidr_block = false
+   cidr_block = "10.0.0.0/16"
+   default_network_acl_id = (known after apply)
+   default_route_table_id = (known after apply)
+   default_security_group_id = (known after apply)
+   dhcp_options_id = (known after apply)
+   enable_classiclink = (known after apply)
+   enable_classiclink_dns_support = (known after apply)
+   enable_dns_hostnames = true
+   enable_dns_support = true
+   enable_network_address_usage_metrics = (known after apply)
+   id = (known after apply)
+   instance_tenancy = "default"
+   ipv6_association_id = (known after apply)
+   ipv6_cidr_block = (known after apply)
+   ipv6_cidr_block_network_border_group = (known after apply)
+   main_route_table_id = (known after apply)
+   owner_id = (known after apply)
+   tags = {
+     + "Name" = "guru-vpc"
+   }
+   tags_all = {
+     + "Name" = "guru-vpc"
+   }
}

# module.eks.module.eks_managed_node_group["one"].aws_eks_node_group.this[0] will be created
+ resource "aws_eks_node_group" "this" {

```

```

+ ami_type           = "AL2_x86_64"
+ arn                = (known after apply)
+ capacity_type      = (known after apply)
+ cluster_name       = (known after apply)
+ disk_size          = (known after apply)
+ id                 = (known after apply)
+ instance_types     = [
  + "t3.small",
]
+ node_group_name     = (known after apply)
+ node_group_name_prefix = "node-group-1-"
+ node_role_arn       = (known after apply)
+ release_version     = (known after apply)
+ resources           = (known after apply)
+ status              = (known after apply)
+ subnet_ids          = (known after apply)
+ tags                = {
  + "Name" = "node-group-1"
}
+ tags_all            = {
  + "Name" = "node-group-1"
}
+ version              = "1.24"

+ launch_template {
  + id      = (known after apply)
  + name    = (known after apply)
  + version = (known after apply)
}

+ scaling_config {
  + desired_size = 2
  + max_size     = 3
  + min_size     = 1
}

+ timeouts {}

+ update_config {
  + max_unavailable_percentage = 33
}
}

# module.eks.module.eks_managed_node_group["one"].aws_iam_role.this[0] will be created
+ resource "aws_iam_role" "this" {
  + arn                = (known after apply)
  + assume_role_policy = jsonencode(
    {
      + Statement = [
        + {
          + Action      = "sts:AssumeRole"
          + Effect      = "Allow"
          + Principal = {
            + Service = "ec2.amazonaws.com"
          }
          + Sid      = "EKSTNodeAssumeRole"
        },
      ]
      + Version = "2012-10-17"
    }
  )
  + create_date      = (known after apply)
  + description      = "EKS managed node group IAM role"
  + force_detach_policies = true
  + id                = (known after apply)
  + managed_policy_arns = (known after apply)
  + max_session_duration = 3600
  + name              = (known after apply)
  + name_prefix       = "node-group-1-eks-node-group-"
  + path              = "/"
  + tags_all          = (known after apply)
  + unique_id         = (known after apply)

  + inline_policy {
    + name = (known after apply)
    + policy = (known after apply)
  }
}

# module.eks.module.eks_managed_node_group["one"].aws_iam_role_policy_attachment.this["arn:aws:iam::aws:policy/AmazonEC2ContainerRegistry"]
+ resource "aws_iam_role_policy_attachment" "this" {

```

```

+ id          = (known after apply)
+ policy_arn  = "arn:aws:iam::aws:policy/AmazonEC2ContainerRegistryReadOnly"
+ role       = (known after apply)
}

# module.eks.module.eks_managed_node_group["one"].aws_iam_role_policy_attachment.this["arn:aws:iam::aws:policy/AmazonEKSWorkerNodePolicy"]
+ resource "aws_iam_role_policy_attachment" "this" {
+ id          = (known after apply)
+ policy_arn  = "arn:aws:iam::aws:policy/AmazonEKSWorkerNodePolicy"
+ role       = (known after apply)
}

# module.eks.module.eks_managed_node_group["one"].aws_iam_role_policy_attachment.this["arn:aws:iam::aws:policy/AmazonEKS_CNI_Policy"] will
+ resource "aws_iam_role_policy_attachment" "this" {
+ id          = (known after apply)
+ policy_arn  = "arn:aws:iam::aws:policy/AmazonEKS_CNI_Policy"
+ role       = (known after apply)
}

# module.eks.module.eks_managed_node_group["one"].aws_launch_template.this[0] will be created
+ resource "aws_launch_template" "this" {
+ arn          = (known after apply)
+ default_version = (known after apply)
+ description   = "Custom launch template for node-group-1 EKS managed node group"
+ id          = (known after apply)
+ latest_version = (known after apply)
+ name        = (known after apply)
+ name_prefix  = "one-"
+ tags_all    = (known after apply)
+ update_default_version = true
+ vpc_security_group_ids = (known after apply)

+ metadata_options {
+ http_endpoint        = "enabled"
+ http_protocol_ipv6   = "disabled"
+ http_put_response_hop_limit = 2
+ http_tokens          = "required"
+ instance_metadata_tags = "disabled"
}

+ monitoring {
+ enabled = true
}

+ tag_specifications {
+ resource_type = "instance"
+ tags         = {
+   "Name" = "node-group-1"
}
}

+ tag_specifications {
+ resource_type = "network-interface"
+ tags         = {
+   "Name" = "node-group-1"
}
}

+ tag_specifications {
+ resource_type = "volume"
+ tags         = {
+   "Name" = "node-group-1"
}
}
}

# module.eks.module.eks_managed_node_group["two"].aws_eks_node_group.this[0] will be created
+ resource "aws_eks_node_group" "this" {
+ ami_type      = "AL2_x86_64"
+ arn           = (known after apply)
+ capacity_type = (known after apply)
+ cluster_name  = (known after apply)
+ disk_size     = (known after apply)
+ id           = (known after apply)
+ instance_types = [
+   "t3.small",
]
+ node_group_name      = (known after apply)
+ node_group_name_prefix = "node-group-2-"
+ node_role_arn        = (known after apply)
+ release_version      = (known after apply)
+ resources            = (known after apply)
+ status              = (known after apply)
}

```

```

+ subnet_ids          = (known after apply)
+ tags                = {
+   + "Name" = "node-group-2"
+ }
+ tags_all            = {
+   + "Name" = "node-group-2"
+ }
+ version              = "1.24"

+ launch_template {
+   + id      = (known after apply)
+   + name    = (known after apply)
+   + version = (known after apply)
+ }

+ scaling_config {
+   + desired_size = 1
+   + max_size     = 2
+   + min_size     = 1
+ }

+ timeouts {}

+ update_config {
+   + max_unavailable_percentage = 33
+ }
}

# module.eks.module.eks_managed_node_group["two"].aws_iam_role.this[0] will be created
+ resource "aws_iam_role" "this" {
+   + arn              = (known after apply)
+   + assume_role_policy = jsonencode(
+     {
+       + Statement = [
+         + {
+           + Action      = "sts:AssumeRole"
+           + Effect      = "Allow"
+           + Principal = {
+             + Service = "ec2.amazonaws.com"
+           }
+           + Sid        = "EKSTNodeAssumeRole"
+         },
+       ]
+       + Version = "2012-10-17"
+     }
+   )
+   + create_date      = (known after apply)
+   + description      = "EKS managed node group IAM role"
+   + force_detach_policies = true
+   + id               = (known after apply)
+   + managed_policy_arns = (known after apply)
+   + max_session_duration = 3600
+   + name             = (known after apply)
+   + name_prefix      = "node-group-2-eks-node-group-"
+   + path             = "/"
+   + tags_all         = (known after apply)
+   + unique_id        = (known after apply)

+   + inline_policy {
+     + name = (known after apply)
+     + policy = (known after apply)
+   }
+ }

# module.eks.module.eks_managed_node_group["two"].aws_iam_role_policy_attachment.this["arn:aws:iam::aws:policy/AmazonEC2ContainerRegistry"]
+ resource "aws_iam_role_policy_attachment" "this" {
+   + id          = (known after apply)
+   + policy_arn  = "arn:aws:iam::aws:policy/AmazonEC2ContainerRegistryReadOnly"
+   + role        = (known after apply)
+ }

# module.eks.module.eks_managed_node_group["two"].aws_iam_role_policy_attachment.this["arn:aws:iam::aws:policy/AmazonEKSWorkerNodePolicy"]
+ resource "aws_iam_role_policy_attachment" "this" {
+   + id          = (known after apply)
+   + policy_arn  = "arn:aws:iam::aws:policy/AmazonEKSWorkerNodePolicy"
+   + role        = (known after apply)
+ }

# module.eks.module.eks_managed_node_group["two"].aws_iam_role_policy_attachment.this["arn:aws:iam::aws:policy/AmazonEKS_CNI_Policy"] wil
+ resource "aws_iam_role_policy_attachment" "this" {
+   + id          = (known after apply)

```

```

+ policy_arn = "arn:aws:iam::aws:policy/AmazonEKS_CNI_Policy"
+ role       = (known after apply)
}

# module.eks.module.eks_managed_node_group["two"].aws_launch_template.this[0] will be created
+ resource "aws_launch_template" "this" {
  + arn                = (known after apply)
  + default_version    = (known after apply)
  + description        = "Custom launch template for node-group-2 EKS managed node group"
  + id                 = (known after apply)
  + latest_version     = (known after apply)
  + name               = (known after apply)
  + name_prefix        = "two-"
  + tags_all           = (known after apply)
  + update_default_version = true
  + vpc_security_group_ids = (known after apply)

  + metadata_options {
    + http_endpoint        = "enabled"
    + http_protocol_ipv6   = "disabled"
    + http_put_response_hop_limit = 2
    + http_tokens          = "required"
    + instance_metadata_tags = "disabled"
  }

  + monitoring {
    + enabled = true
  }

  + tag_specifications {
    + resource_type = "instance"
    + tags          = {
      + "Name" = "node-group-2"
    }
  }

  + tag_specifications {
    + resource_type = "network-interface"
    + tags          = {
      + "Name" = "node-group-2"
    }
  }

  + tag_specifications {
    + resource_type = "volume"
    + tags          = {
      + "Name" = "node-group-2"
    }
  }
}

# module.eks.module.kms.data.aws_iam_policy_document.this[0] will be read during apply #
(config refers to values not yet known)
<= data "aws_iam_policy_document" "this" {
  + id          = (known after apply)
  + json        = (known after apply)
  + override_policy_documents = []
  + source_policy_documents = []

  + statement {
    + actions = [
      + "kms:CancelKeyDeletion",
      + "kms:Create*",
      + "kms:Delete*",
      + "kms:Describe*",
      + "kms:Disable*",
      + "kms:Enable*",
      + "kms:Get*",
      + "kms:List*",
      + "kms:Put*",
      + "kms:Revoke*",
      + "kms:ScheduleKeyDeletion",
      + "kms:TagResource",
      + "kms:UntagResource",
      + "kms:Update*",
    ]
    + resources = [
      + "*",
    ]
    + sid       = "KeyAdministration"

    + principals {
      + identifiers = [

```

```

        + "arn:aws:iam::988505337009:user/k8sadmin",
      ]
      + type      = "AWS"
    }
  }
+ statement {
  + actions = [
    + "kms:Decrypt",
    + "kms:DescribeKey",
    + "kms:Encrypt",
    + "kms:GenerateDataKey*",
    + "kms:ReEncrypt*",
  ]
  + resources = [
    + "*",
  ]
  + sid      = "KeyUsage"

  + principals {
    + identifiers = [
      + (known after apply),
    ]
    + type      = "AWS"
  }
}

# module.eks.module.kms.aws_kms_alias.this["cluster"] will be created
+ resource "aws_kms_alias" "this" {
  + arn          = (known after apply)
  + id           = (known after apply)
  + name         = (known after apply)
  + name_prefix  = (known after apply)
  + target_key_arn = (known after apply)
  + target_key_id = (known after apply)
}

# module.eks.module.kms.aws_kms_key.this[0] will be created
+ resource "aws_kms_key" "this" {
  + arn          = (known after apply)
  + bypass_policy_lockout_safety_check = false
  + customer_master_key_spec          = "SYMMETRIC_DEFAULT"
  + description                       = (known after apply)
  + enable_key_rotation               = true
  + id                               = (known after apply)
  + is_enabled                        = true
  + key_id                           = (known after apply)
  + key_usage                         = "ENCRYPT_DECRYPT"
  + multi_region                     = false
  + policy                           = (known after apply)
  + tags_all                         = (known after apply)
}
Plan: 54 to add, 0 to change, 0 to destroy.
Changes to Outputs:
  + cluster_endpoint      = (known after apply)
  + cluster_name           = (known after apply)
  + cluster_security_group_id = (known after apply)
  + region                 = "us-east-1"

```

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

Finalmente terraform apply (Lo cual nos va a llevar unos cuantos minutos).

```
terraform apply --auto-approve
```

Finalmente

```
Apply complete! Resources: 54 added, 0 changed, 0 destroyed.

Outputs:

cluster_endpoint = "https://43f725ef95a10f6ff084d18a8b0fa2bd.gr7.us-east-1.eks.amazonaws.com"
cluster_name     = "guru-eks-eks1c08n"
cluster_security_group_id =
"sg-0b3b25f18a24fdf0d" region = "us-east-1"
```

ahora agregamos el context

```
aws eks --region ${terraform output -raw region} update-kubeconfig --name ${terraform output -raw cluster_name}
```

vemos que está todo ok, seguimos

```
[ec2-user@ip-172-31-30-154 eks]$ aws eks --region ${terraform output -raw region} update-kubeconfig --name ${terraform output -raw cluster_name}
Added new context arn:aws:eks:us-east-1:988505337009:cluster/guru-eks-eks1c08n to /home/ec2-user/.kube/config
[ec2-user@ip-172-31-30-154 eks]$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-10-0-1-196.ec2.internal	Ready	<none>	5m31s	v1.24.10-eks-48e63af
ip-10-0-2-128.ec2.internal	Ready	<none>	5m10s	v1.24.10-eks-48e63af
ip-10-0-2-135.ec2.internal	Ready	<none>	5m21s	v1.24.10-eks-48e63af

Bajamos la aplicacion del pacman

```
cd ..
[ec2-user@ip-172-31-30-154 ~]$ wget https://github.com/pluralsight-cloud/content-deploying-and-managing-a-web-application-in-kubernetes-with-terraform
--2023-03-08 23:39:51-- https://github.com/pluralsight-cloud/content-deploying-and-managing-a-web-application-in-kubernetes-with-terraform
Resolving github.com (github.com)... 140.82.113.3
Connecting to github.com (github.com)|140.82.113.3|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://raw.githubusercontent.com/pluralsight-cloud/content-deploying-and-managing-a-web-application-in-kubernetes-with-terraform
--2023-03-08 23:39:51-- https://raw.githubusercontent.com/pluralsight-cloud/content-deploying-and-managing-a-web-application-in-kubernetes
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, ... Connecting
to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200
OK Length: 4195 (4.1K) [application/zip]
Saving to: 'pac-man.zip'

100%[----->] 4.195  --K/s in
0s

2023-03-08 23:39:51 (42.3 MB/s) - 'pac-man.zip' saved
(4195/4195)
```

ahora descomprimo

```
[ec2-user@ip-172-31-30-154 ~]$ unzip
pac-man.zip Archive: pac-man.zip
  creating: pac-man/modules/
  inflating: pac-man/pac-man.tf
  creating: pac-man/modules/mongo/
  creating: pac-man/modules/pac-man/
  inflating: pac-man/modules/mongo/mongo-deployment.tf
  inflating: pac-man/modules/mongo/mongo-service.tf
  inflating: pac-man/modules/mongo/mongo-pvc.tf
  inflating: pac-man/modules/mongo/mongo-svc.tf
```



```

inflating:
pac-man/modules/mongo/mongo-service.tf
extracting: pac-man/modules/mongo/variables.tf
inflating:
pac-man/modules/pac-man/pac-man-deployment.tf
inflating: pac-man/modules/pac-man/pac-man-service.tf
extracting: pac-man/modules/pac-man/variables.tf

```

Veo los directorios:

```

[ec2-user@ip-172-31-30-154 ~]$ cd
pac-man/ [ec2-user@ip-172-31-30-154
pac-man]$ ls -R
.:
modules pac-man.tf

./modules:
mongo pac-man

./modules/mongo:
mongo-deployment.tf mongo-pvc.tf mongo-pv.tf mongo-sc.tf mongo-service.tf variables.tf

./modules/pac-man:
pac-man-deployment.tf pac-man-service.tf variables.tf

```

tenemos que agregar la imagen de docker de pacman. vamos a hacerlo!

Línea 30, agregamos.

```

sp
ec
{
  container {
    name = "pac-man"
    image =
      "docker.io/jesasehoch/pacman-nodejs-app:latest"
  }
}

```

Salvamos y vamos al módulo de mongodb

Editamos el archivo

```

[ec2-user@ip-172-31-30-154 pac-man]$ vim pac-man.tf

```

agregamos al final de todo esto

en kubernetes_namespace le faltaba el "pac-man"

```

module "mongo" {
  source = "../modules/mongo"
  kubernetes_namespace = "pac-man"
}

module "pac-man" {
  source = "../modules/pac-man"
  kubernetes_namespace =
    "pac-man" depends_on =
    [module.mongo]
}

```

Y salvamos

```
terraform init && terraform plan && terraform apply
```

Esperamos y vemos si todo salio como esperabamos

```
kubectl -n pac-man get all
```

Vemos que la salida es....

```

NAME                                READY    STATUS    RESTARTS   AGE
pod/mongo-695f6597f4-xkwn9         1/1      Running   0           84s
pod/pac-man-5/955c9b46-kxmtb       1/1      Running   0           68s

NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)
service/mongo       LoadBalancer  172.20.109.60   a5bfda21f0ab34673a36d8ddf525d2d-1856496249.us-east-1.elb.amazonaws.com  27017:31214/TCP
service/pac-man     LoadBalancer  172.20.16.84    aab5384c832ba4ab9abe3e93/994eb4a-1668952125.us-east-1.elb.amazonaws.com  80:32504/TCP

NAME                READY    UP-TO-DATE    AVAILABLE    AGE
deployment.apps/mongo  1/1      1              1            84s
deployment.apps/pac-man  1/1      1              1            68s

NAME                DESIRED    CURRENT    READY    AGE
replicaset.apps/mongo-695f6597f4  1          1          1        84s
replicaset.apps/pac-man-5/955c9b46  1          1          1        68s

```

Lo que nos interesa es

LB de aca. veamos....

```
kubectl get svc -n pacman
```